Peer Review: Team19 Demographic Group Bias in Brain stroke Data

Liujia Yu, Yun-Jing Chang

(Dated: April 9, 2025)

Summary

The purpose of the study is to develop a machine learning model for stroke risk prediction while addressing potential demographic biases in the dataset. The study investigates how demographic factors such as age, gender, and employment type might affect model predictions and aims to ensure fairness across demographic groups. The methodology involves using machine learning algorithms, such as decision trees and random forests, to predict stroke risk based on demographic, health-related, and behavioral data. The model's performance is evaluated using accuracy, precision, recall, and F1-score, while fairness is assessed using statistical parity and disparate impact. Bias mitigation techniques like re-weighting and adversarial debiasing are employed to ensure the model is fair across demographic groups. The study has made progress in developing a machine learning model for stroke risk prediction while addressing demographic biases. Key findings include how factors such as age, gender, and employment type influence the accuracy and fairness of predictions. The project has used bias mitigation techniques (e.g., re-weighting and adversarial debiasing) to reduce disparities in the model's predictions, aiming to create a fair, reliable, and unbiased stroke risk prediction model.

Strengths

1. **Comprehensive approach to bias detection:** The work effectively addresses demographic biases in the stroke risk prediction model, including factors like gender, age, and employment type. By using fairness metrics such as statistical parity and disparate impact, the model ensures it's accurate and fair across demographic groups. This approach to detecting and reducing bias sets the study apart from other stroke prediction models that may overlook fairness.

2. **Clear methodology and evaluation:** The methodology is well-structured and easy to follow, using machine learning algorithms like decision trees and random forests to predict stroke risk. The model is evaluated using metrics like accuracy, precision, recall, and F1-score, as well as fairness metrics. This dual focus on accuracy and fairness ensures the model is reliable and robust.

3. **Use of bias mitigation techniques:** The study doesn't just identify biases but actively works to reduce them with techniques like reweighting and adversarial debiasing. These methods improve fairness in the model's predictions, making it more suitable for real-world healthcare applications. This shows a commitment to ensuring AI-driven solutions are fair and don't unfairly favor certain demographic groups.

Weaknesses

1. **Limited exploration of alternative bias mitigation techniques:** While the study uses re-weighting and adversarial debiasing, it could benefit from exploring a broader range of bias mitigation strategies. Techniques like causal fairness, fairness-aware learning, or post-processing methods could offer improvements in ensuring fairness across diverse demographic groups. Expanding the scope of bias mitigation strategies could help fill gaps and make the model more robust.

2. **Potential imbalance in data representation:** Even though some biases are addressed, the dataset may still have issues with underrepresentation or overrepresentation of certain demographic groups, particularly in terms of gender and work type. This imbalance could affect the model's ability to make accurate predictions for underrepresented groups. A deeper dive into the dataset's distribution, along with techniques like SMOTE or resampling methods, could help address this and improve fairness.

3. **Limited evaluation of real-world application:** While the study evaluates fairness within the dataset, it doesn't fully assess how well the model performs in real-world scenarios, especially in diverse healthcare settings. To make the model robust, it would be important to validate it using external datasets or through cross-validation with real-world healthcare data. This would ensure that the model works equitably outside the controlled experimental environment.

Other Comments

In terms of visualizations, confusion matrices or ROC curves could provide additional insights into the model's performance. This would be useful for evaluating class imbalances and predictive accuracy. Additionally, it might be worth considering grouping "Work Type" categories (e.g., "Private" and "Self-employed") if the dataset is small or if the differences between categories don't impact the model's performance. This could reduce complexity and make the model more generalizable.